


Copyright Notice

These slides are distributed under the Creative Commons License.

[DeepLearning.AI](#) makes these slides available for educational purposes. You may not use or distribute these slides for commercial purposes. You may make copies of these slides and use or distribute them for educational purposes as long as you cite [DeepLearning.AI](#) as the source of the slides.

For the rest of the details of the license, see <https://creativecommons.org/licenses/by-sa/2.0/legalcode>

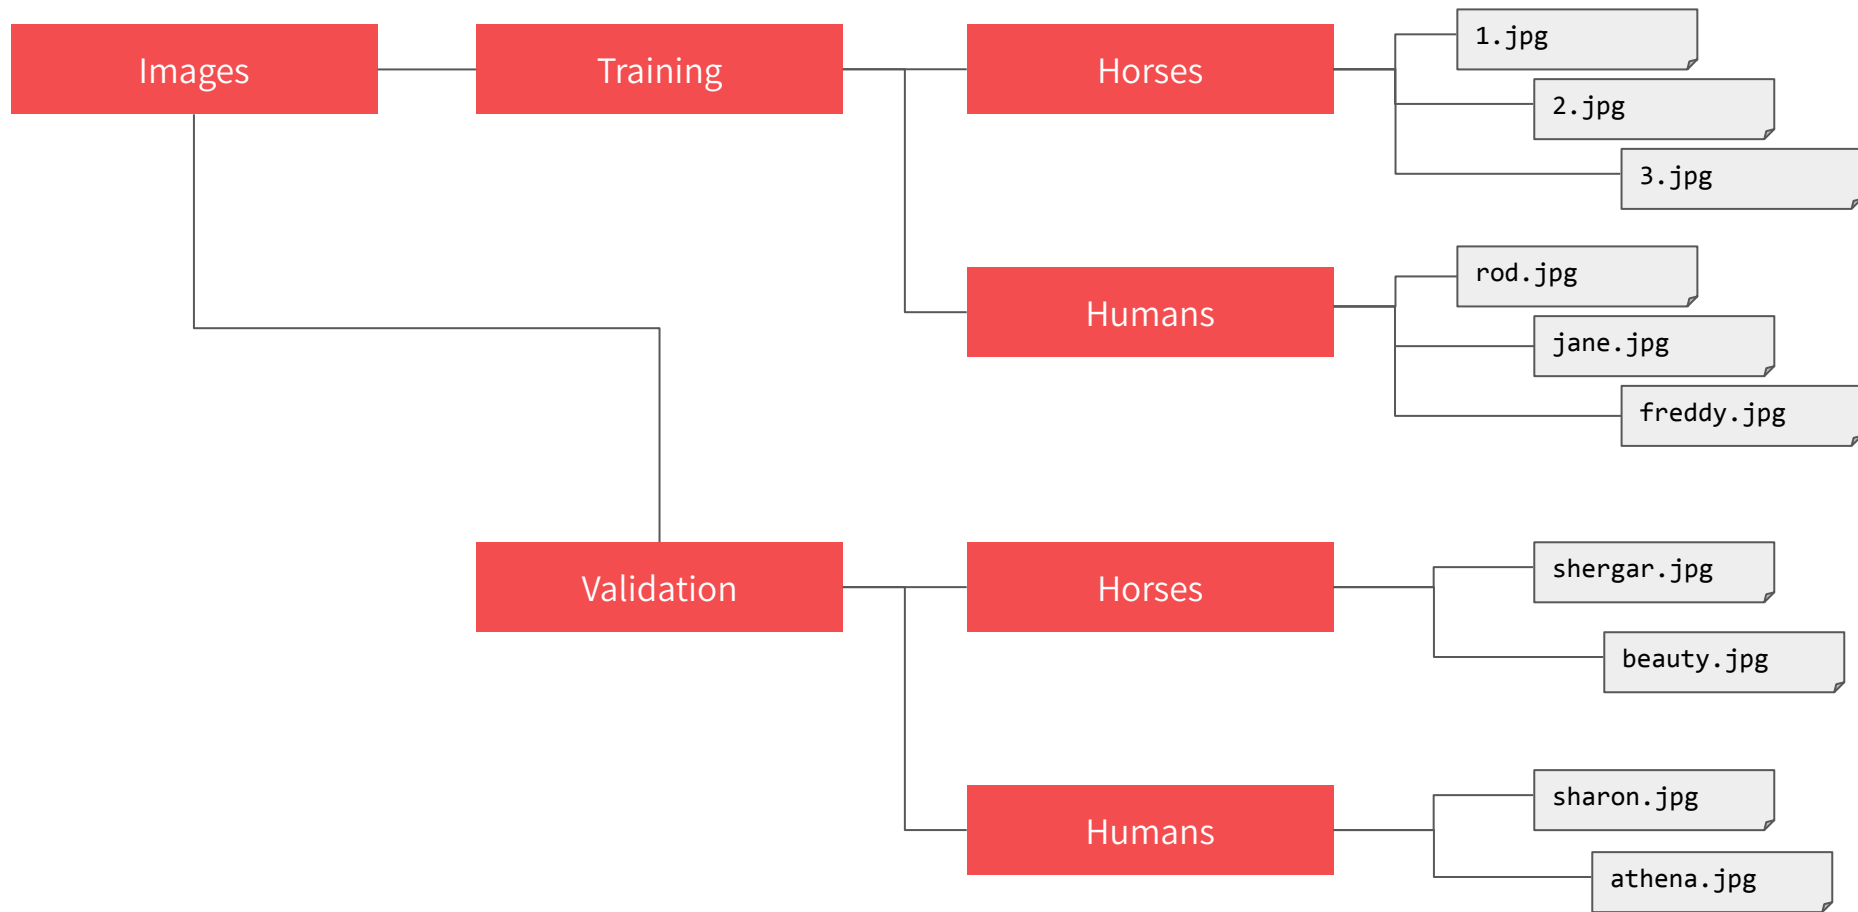



tf.data API

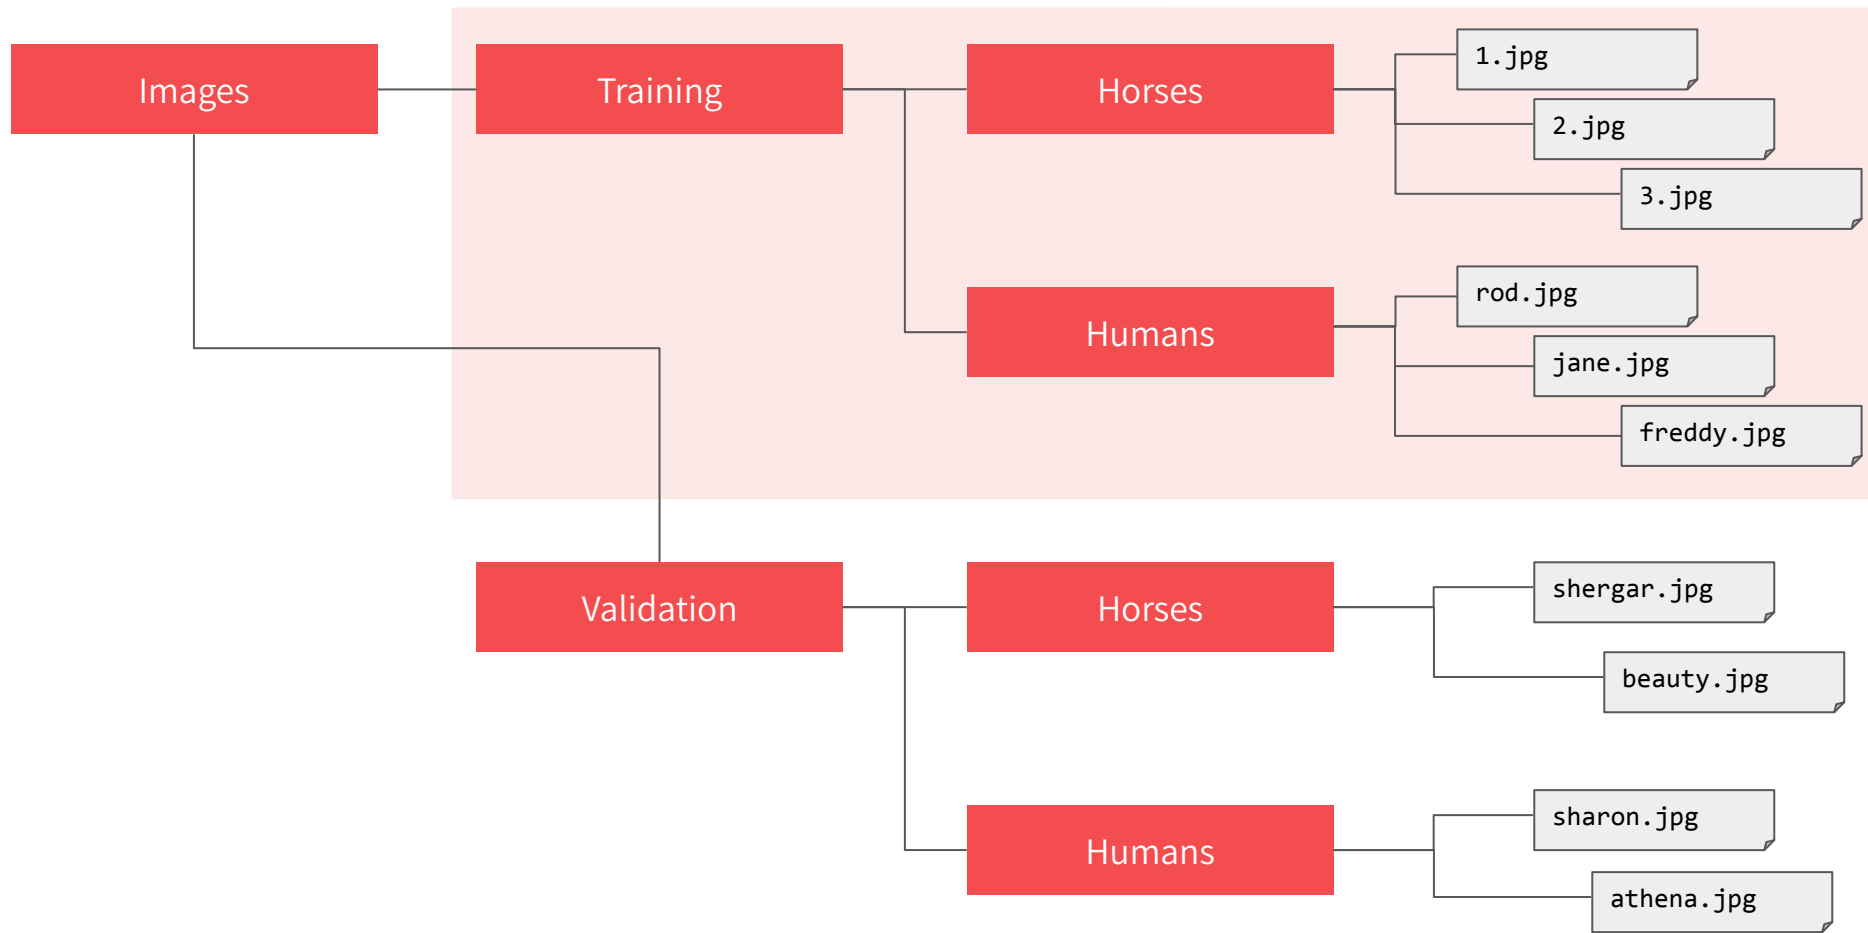


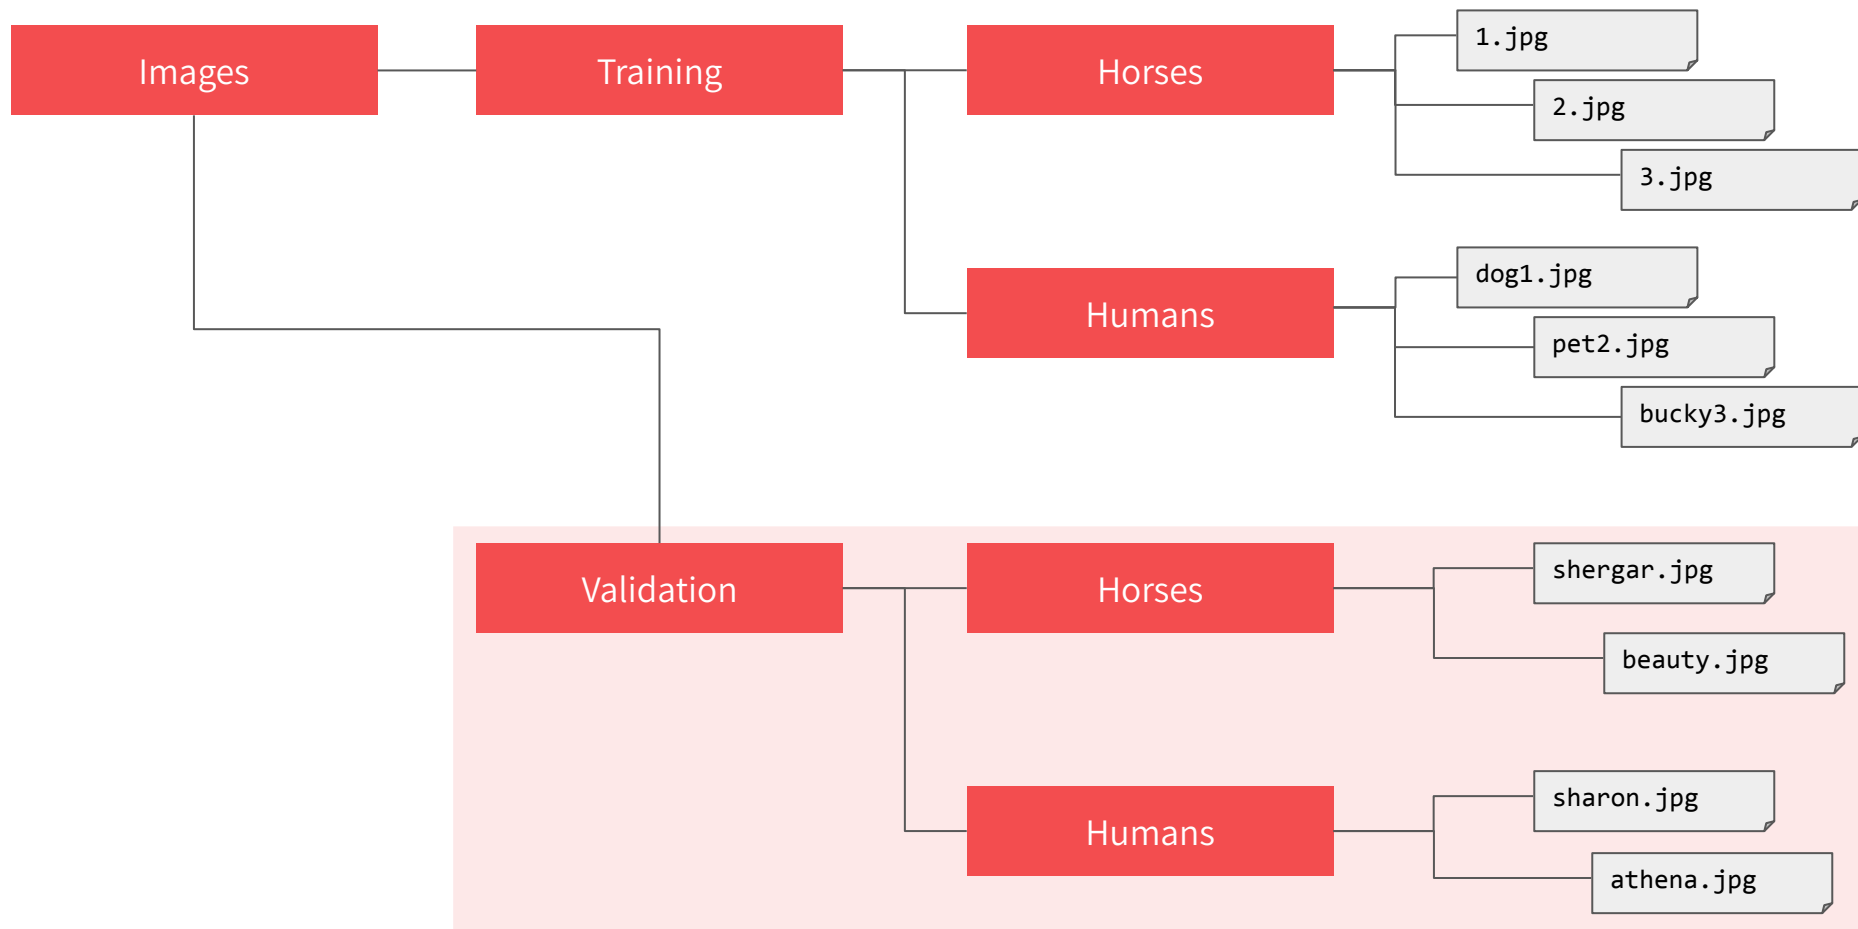
tf.keras.utils.image_dataset_from_directory












```
train_dataset = tf.keras.utils.image_dataset_from_directory(  
    TRAIN_DIR,  
    image_size=(300, 300),  
    batch_size=128,  
    label_mode='binary'  
)
```



```
train_dataset = tf.keras.utils.image_dataset_from_directory(  
    TRAIN_DIR,  
    image_size=(300, 300),  
    batch_size=128,  
    label_mode='binary'  
)
```



```
train_dataset = tf.keras.utils.image_dataset_from_directory(  
    TRAIN_DIR,  
    image_size=(300, 300),  
    batch_size=128,  
    label_mode='binary'  
)
```



```
train_dataset = tf.keras.utils.image_dataset_from_directory(  
    TRAIN_DIR,  
    image_size=(300, 300),  
    batch_size=128,  
    label_mode='binary'  
)
```



tf.data.Dataset

https://www.tensorflow.org/api_docs/python/tf/data/Dataset

```
rescale_layer = tf.keras.layers.Rescaling(scale=1./255)
```

```
train_dataset_scaled = train_dataset.map(  
    lambda image, label: (rescale_layer(image), label))
```



```
rescale_layer = tf.keras.layers.Rescaling(scale=1./255)

train_dataset_scaled = train_dataset.map(
    lambda image, label: (rescale_layer(image), label))
```



```
rescale_layer = tf.keras.layers.Rescaling(scale=1./255)

train_dataset_scaled = train_dataset.map(
    lambda image, label: (rescale_layer(image), label))
```




```
train_dataset_final = (train_dataset_scaled  
                        .cache()  
                        .shuffle(buffer_size=1000)  
                        .prefetch(buffer_size=tf.data.AUTOTUNE))
```

```
train_dataset_final = (train_dataset_scaled  
                        .cache()  
                        .shuffle(buffer_size=1000)  
                        .prefetch(buffer_size=tf.data.AUTOTUNE))
```

```
train_dataset_final = (train_dataset_scaled  
                        .cache()  
                        .shuffle(buffer_size=1000)  
                        .prefetch(buffer_size=tf.data.AUTOTUNE))
```

```
train_dataset_final = (train_dataset_scaled  
                        .cache()  
                        .shuffle(buffer_size=1000)  
                        .prefetch(buffer_size=tf.data.AUTOTUNE))
```



```
validation_dataset = tf.keras.utils.image_dataset_from_directory(
    VAL_DIR,
    image_size=(300, 300),
    batch_size=32,
    label_mode='binary'
)

validation_dataset_scaled = validation_dataset.map(lambda image, label:
(rescale_layer(image), label))

# Configure the validation dataset
validation_dataset_final = (validation_dataset_scaled
    .cache()
    .prefetch(buffer_size=tf.data.AUTOTUNE))
```



```
validation_dataset = tf.keras.utils.image_dataset_from_directory(
    VAL_DIR,
    image_size=(300, 300),
    batch_size=128,
    label_mode='binary'
)
```

```
validation_dataset_scaled = validation_dataset.map(lambda image, label:
    (rescale_layer(image), label))
```

```
# Configure the validation dataset
```

```
validation_dataset_final = (validation_dataset_scaled
    .cache()
    .prefetch(buffer_size=tf.data.AUTOTUNE))
```



```
model = tf.keras.models.Sequential([
    tf.keras.layers.Input(shape=(300, 300, 3)),
    tf.keras.layers.Conv2D(16, (3, 3), activation='relu'),
    tf.keras.layers.MaxPooling2D(2, 2),
    tf.keras.layers.Conv2D(32, (3, 3), activation='relu'),
    tf.keras.layers.MaxPooling2D(2, 2),
    tf.keras.layers.Conv2D(64, (3, 3), activation='relu'),
    tf.keras.layers.MaxPooling2D(2, 2),
    tf.keras.layers.Flatten(),
    tf.keras.layers.Dense(512, activation='relu'),
    tf.keras.layers.Dense(1, activation='sigmoid')
])
```

```
model = tf.keras.models.Sequential([
    tf.keras.layers.Input(shape=(300, 300, 3)),
    tf.keras.layers.Conv2D(16, (3, 3), activation='relu'),
    tf.keras.layers.MaxPooling2D(2, 2),
    tf.keras.layers.Conv2D(32, (3, 3), activation='relu'),
    tf.keras.layers.MaxPooling2D(2, 2),
    tf.keras.layers.Conv2D(64, (3, 3), activation='relu'),
    tf.keras.layers.MaxPooling2D(2, 2),
    tf.keras.layers.Flatten(),
    tf.keras.layers.Dense(512, activation='relu'),
    tf.keras.layers.Dense(1, activation='sigmoid')
])
```



```
model = tf.keras.models.Sequential([
    tf.keras.layers.Input(shape=(300, 300, 3)),
    tf.keras.layers.Conv2D(16, (3, 3), activation='relu'),
    tf.keras.layers.MaxPooling2D(2, 2),
    tf.keras.layers.Conv2D(32, (3, 3), activation='relu'),
    tf.keras.layers.MaxPooling2D(2, 2),
    tf.keras.layers.Conv2D(64, (3, 3), activation='relu'),
    tf.keras.layers.MaxPooling2D(2, 2),
    tf.keras.layers.Flatten(),
    tf.keras.layers.Dense(512, activation='relu'),
    tf.keras.layers.Dense(1, activation='sigmoid')
])
```

```
model = tf.keras.models.Sequential([
    tf.keras.layers.Input(shape=(300, 300, 3)),
    tf.keras.layers.Conv2D(16, (3, 3), activation='relu'),
    tf.keras.layers.MaxPooling2D(2, 2),
    tf.keras.layers.Conv2D(32, (3, 3), activation='relu'),
    tf.keras.layers.MaxPooling2D(2, 2),
    tf.keras.layers.Conv2D(64, (3, 3), activation='relu'),
    tf.keras.layers.MaxPooling2D(2, 2),
    tf.keras.layers.Flatten(),
    tf.keras.layers.Dense(512, activation='relu'),
    tf.keras.layers.Dense(1, activation='sigmoid')
])
```

Layer (type)	Output Shape	Param #
conv2d_5 (Conv2D)	(None, 298, 298, 16)	448
max_pooling2d_5 (MaxPooling2D)	(None, 149, 149, 16)	0
conv2d_6 (Conv2D)	(None, 147, 147, 32)	4640
max_pooling2d_6 (MaxPooling2D)	(None, 73, 73, 32)	0
conv2d_7 (Conv2D)	(None, 71, 71, 64)	18496
max_pooling2d_7 (MaxPooling2D)	(None, 35, 35, 64)	0
flatten_1 (Flatten)	(None, 78400)	0
dense_2 (Dense)	(None, 512)	40141312
dense_3 (Dense)	(None, 1)	513
Total params: 40,165,409		
Trainable params: 40,165,409		
Non-trainable params: 0		

```
model.compile(loss='binary_crossentropy',  
              optimizer=tf.keras.optimizers.RMSprop(learning_rate=0.001),  
              metrics=['accuracy'])
```

<https://youtu.be/zLRB4oupj6g>



Machine Learning



0:06 / 8:58



2.1.4 Gradient Descent in Practice II Learning Rate by Andrew Ng

```
history = model.fit(  
    train_dataset_final,  
    epochs=15,  
    validation_data=validation_dataset_final,  
    verbose=2)
```



```
history = model.fit(  
    train_dataset_final,  
    epochs=15,  
    validation_data=validation_dataset_final,  
    verbose=2)
```



```
history = model.fit(  
    train_dataset_final,  
    epochs=15,  
    validation_data=validation_dataset_final,  
    verbose=2)
```



```
history = model.fit(  
    train_dataset_final,  
    epochs=15,  
    validation_data=validation_dataset_final,  
    verbose=2)
```



```
history = model.fit(  
    train_dataset_final,  
    epochs=15,  
    validation_data=validation_dataset_final,  
    verbose=2)
```



```
from google.colab import files

uploaded = files.upload()

for filename in uploaded.keys():

    # predicting images
    path = '/content/' + filename
    image = tf.keras.utils.load_img(path, target_size=(300, 300))
    image = tf.keras.utils.img_to_array(image)
    image = rescale_layer(image)
    image = np.expand_dims(image, axis=0)

    prediction = model.predict(image, verbose=0)[0][0]
    print(f'\nmodel output: {prediction}')

    if prediction > 0.5:
        print(filename + " is a human")
    else:
        print(filename + " is a horse")
```



```
from google.colab import files
```

```
uploaded = files.upload()
```

```
for filename in uploaded.keys():
```

```
    # predicting images
```

```
    path = '/content/' + filename
```

```
    image = tf.keras.utils.load_img(path, target_size=(300, 300))
```

```
    image = tf.keras.utils.img_to_array(image)
```

```
    image = rescale_layer(image)
```

```
    image = np.expand_dims(image, axis=0)
```

```
    prediction = model.predict(image, verbose=0)[0][0]
```

```
    print(f'\nmodel output: {prediction}')
```

```
    if prediction > 0.5:
```

```
        print(filename + " is a human")
```

```
    else:
```

```
        print(filename + " is a horse")
```



```
from google.colab import files
```

```
uploaded = files.upload()
```

```
for filename in uploaded.keys():
```

```
    # predicting images
```

```
    path = '/content/' + filename
```

```
    image = tf.keras.utils.load_img(path, target_size=(300, 300))
```

```
    image = tf.keras.utils.img_to_array(image)
```

```
    image = rescale_layer(image)
```

```
    image = np.expand_dims(image, axis=0)
```

```
    prediction = model.predict(image, verbose=0)[0][0]
```

```
    print(f'\nmodel output: {prediction}')
```

```
    if prediction > 0.5:
```

```
        print(filename + " is a human")
```

```
    else:
```

```
        print(filename + " is a horse")
```



```
from google.colab import files

uploaded = files.upload()

for filename in uploaded.keys():

    # predicting images
    path = '/content/' + filename
    image = tf.keras.utils.load_img(path, target_size=(300, 300))
    image = tf.keras.utils.img_to_array(image)
    image = rescale_layer(image)
    image = np.expand_dims(image, axis=0)

    prediction = model.predict(image, verbose=0)[0][0]
    print(f'\nmodel output: {prediction}')

    if prediction > 0.5:
        print(filename + " is a human")
    else:
        print(filename + " is a horse")
```

